

# Informatics 1: Data & Analysis

## Lecture 3: The Relational Model

Ian Stark

School of Informatics  
The University of Edinburgh

Tuesday 22 January 2019  
Semester 2 Week 2

# Lecture Plan for Weeks 1–4

## Data Representation

This first course section starts by presenting two common **data representation models**.

- The *entity-relationship (ER)* model
- The *relational* model

Note slightly different naming:  
-relationship vs. relational

## Data Manipulation

This is followed by some methods for manipulating data in the relational model and using it to extract information.

- *Relational algebra*
- The *tuple-relational calculus*
- The query language *SQL*


# Outline

- 1 Review of Entities and Relationships
- 2 Constraints on Relationships
- 3 Basics of the Relational Model

# Entity-Relationship Modelling

## The Story so Far

- Requirements Analysis, Conceptual Design, Logical Design
- Entities: Entity Instances, Entity Sets; Attributes, Domains
- Keys: Superkeys, Candidate Keys, Primary Keys, Composite Keys
- Relationships: Relationship Instances, Relationship Sets ← I'll do that one again



# Relationships

A *relationship* is an association between entities. For example, the **takes** relationship between students and courses.


Each individual occurrence of the relationship is a *relationship instance*, and the collection of all such is a *relationship set*.

Relationship	Takes
Relationship instance	((s0456782), (INF08013, 2018))
Relationship set	Edinburgh course registrations

# Relationships


ER diagrams show relationships as diamonds, labelled with the name of the relationship and connected to all the participating entities.

A relationship may also have its own attributes.



# Relationships

Relationships can be between two entities (“**binary**”), three (“**ternary**”) or more (“**n-ary**”).

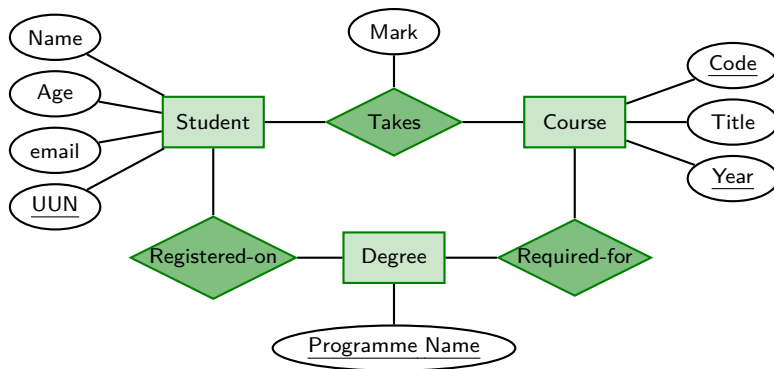


**Relationship instance** (ASDA, Heinz, Ketchup, Edinburgh, £2)

# Relationships

There is no bound on the number of entities participating in a relationship.

An entity may be involved in any number of different relationships.





# Summary

The **entity-relationship** (ER) model is a way to organise the description of *entities* (individual things) and the *relationships* between them.

So far we have seen the following elements of ER modelling:

- **Entities**, all with characteristic **attributes**;
- For each kind of entity there are **entity instances** that together make up an **entity set**;
- A **key** as a minimal set of attributes to identify and distinguish entity instances;
- **Relationships** between entities, which may be **binary**, **ternary**, or **n-ary**.

Entity-relationship modelling provides a *graphical language* for describing this structure in an **ER diagram**.


# Outline

- 1 Review of Entities and Relationships
- 2 Constraints on Relationships
- 3 Basics of the Relational Model

# Constraints on Relationships

In general, relationships need have no distinguished direction between the entities involved, nor limits on how many entity instances are related.

However, some relationships are more constrained.




# Constraints on Relationships

A binary relationship  $R$  between entities  $A$  and  $B$  can be:

**Many-to-one:** Several  $A$  may relate to a single  $B$ ; but not the other way round.

**One-to-many:** Each  $A$  may relate to several  $B$ ; but not the other way round.

**Many-to-many:** Any number of  $A$  may be related to any number of  $B$ .




## Examples?

Container **On** Ship; Artist **Paints** Portrait; Student **Takes** Course.

# Key Constraints

An entity  $E$  has a *key constraint* in relationship  $R$  if each entity instance  $x$  of  $E$  can appear in at most one relationship instance  $(x, y, z, \dots)$  from  $R$ .


An ER diagram indicates a key constraint with an arrowhead on the line joining entity set  $E$  to relationship set  $R$ .



# Key Constraints

## Note

There is a subtlety here: a key constraint does not just report that entity instances happen to appear in no more than one relationship instance — instead, it asserts a structural property of the data that each entity instance **cannot** be in more than one relationship instance.



# Participation Constraints

**Total Participation** of entity E in relationship R means that every entity instance  $x$  in E appears in at least one relationship instance of R.

**Partial Participation** of entity E in relationship R means that some entity instances in E might appear in no relationship instance from R.


## Examples?

Student **Speaks** Language; Country **Uses** Currency; Course **Has** Exam

Again this indicates not just chance but a structural constraint on the data model: each entity instance **must** appear in the relationship.

# Participation Constraints

ER diagrams indicate total participation with a thick or double line from the entity to the relationship.



Again this indicates not just chance but a structural constraint on the data model: each entity instance **must** appear in the relationship.




# Weak Entities

Sometimes the attributes of an entity are not enough to build a key that uniquely identifies each individual.

We may be able to fix this using attributes from other, related, entities.

For this the entity must have a **key constraint** and **total participation** in some *identifying relationship*.


It is then a *weak entity*, with borrowed attributes coming from the *identifying owner*.



# Weak Entities

ER diagrams represent **weak entities** by:

- A double or thick border on the rectangle of the weak entity;
- A double or thick border on the diamond of the identifying relationship;
- A double or thick line and arrow linking these;
- A double or dashed underline for the attributes of the weak entity that contribute to the composite key.

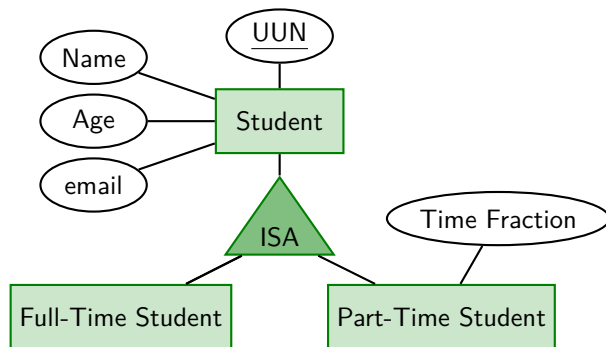


# Entity Hierarchies and Inheritance

Sometimes one entity will refine another: a *subclass* entity *specializes* a *superclass* entity.

Each subclass entity *inherits* the attributes of the superclass entity, and may add its own attributes, too.

ER diagrams represent inheritance with an *IS-A* triangle.



# Entity-Relationship Refinements

## What We've Just Seen

- Relationships that are: One-to-Many, Many-to-One, Many-to-Many
- Key constraints
- Participation constraints: Total participation, Partial participation
- Weak entities with an identifying relationship and identifying owner
- Entity hierarchies and inheritance

**Whitehead:** Civilization advances by extending the number of important operations which we can perform without thinking about them.

# Homework from Tuesday

## Read This

Sections 2.1–2.4 of this textbook on databases: you will need to log in to *Learn* to access this chapter as a PDF.



R. Ramakrishnan and J. Gehrke.  
*Database Management Systems*.  
McGraw-Hill, third edition, 2003.

This is the recommended textbook for the third-year course **Database Systems**. It's a large book, with thorough and extensive material on a wide range of database topics.

It is *not* necessary to buy this book for Inf1-DA. Instead:

## Do This

Step inside the library, locate the HUB and find this book. Have a look at the other textbooks there, too, and compare style and content.

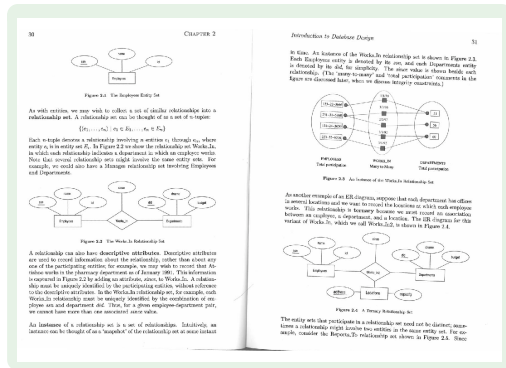
**Finally:** Find the video of *this lecture* online and watch the first minute. If you have any problems then ask on Piazza/Facebook.

# Homework

## Read This

Before the next lecture, on Friday, read the remaining sections, §§2.5 onwards, of Ramakrishnan and Gehrke, completing Chapter 2.

These consider trade-offs and choices in the design of Entity-Relationship models, as well as more on the wider context of modelling.



R. Ramakrishnan and J. Gehrke.  
*Database Management Systems*.  
McGraw-Hill, third edition, 2003.



- Be there at every lecture
- Take notes, paper or electronic (slides will be available online and as paper handouts)
- Write up your notes after the lecture: rearrange, reflect, summarise
- Do the homework: read this, do that, watch these
- Work through the exercises before each tutorial
- Be there at every tutorial: take part, work together with others
- After each tutorial, do some of the additional example questions
- If you have questions or problems then ask for help: on Piazza or Facebook; other students, tutors, me; in person, by email, after lectures.



Tutorials start in **Week 3** of semester and continue each week until the end of semester, except for during the *Festival of Creative Learning* week which falls between lecturing weeks 5 and 6.

Your *course tutor* leads your tutorial group; this is not the same as your *personal tutor*.


You will be able to choose your own tutorial group during Week 2: all tutorials will be one-hour time-slots on Monday, Tuesday or Wednesday.

Studying for Inf-DA requires you to work through exercise sheets and then take part in your weekly tutorials. If you are ill or otherwise unable to attend one week then email your course tutor, and if possible attend another tutorial group in the same week.

# Outline

- 1 Review of Entities and Relationships
- 2 Constraints on Relationships
- 3 Basics of the Relational Model

## A DATABASE SYSTEM



The *relational model* was introduced in 1970 by Edgar F. Codd, a British computer scientist working for IBM in California.

A key insight, and a reason for the success of the relational model, is its separation between *specification* (what you want to find out) and *implementation* (how the system should find it out for you).


IBM were initially reluctant to exploit Codd's idea, but did in due course develop the **System R** database platform which led the commercial development of *relational database management systems* (RDBMS).



System R included the **SEQUEL** language, which then became **SQL** and the standard query language for all subsequent relational databases.

Relational database management systems are now a very large industry: G\$38 market size in 2014, and growing.

Codd received the **1981 Turing Award** for his work on databases.

## Global Database Market (\$B)

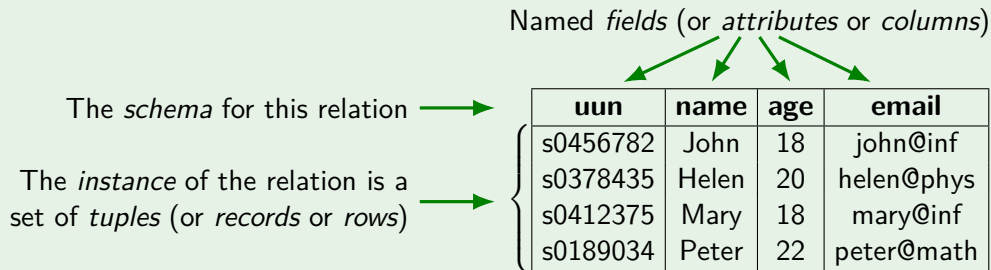




Credits: IBM, Ferkel

# The Relational Model

Relational databases take as fundamental the idea of a *relation*, comprising a *schema* and an *instance*.



**Absolutely everything** in a relational database is built from relations and operations upon them.

# The Relational Model

Relational databases take as fundamental the idea of a *relation*, comprising a *schema* and an *instance*.

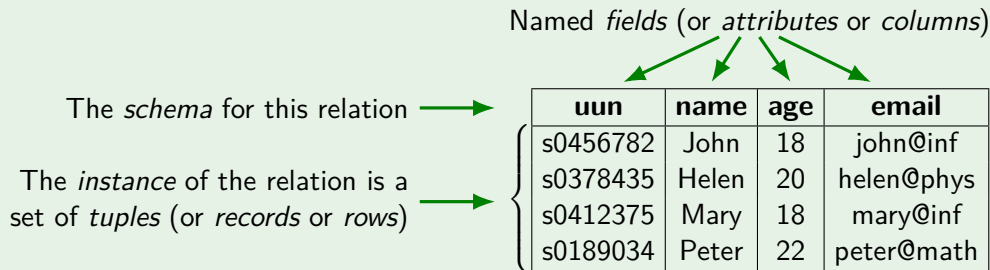
- The **schema** is the format of the relation:
  - A set of named *fields* (or *attributes* or *columns*)
  - For each field its *domain* (or *type*)
- The **instance** of a relation is a *table*:
  - A set of *rows* (or *records* or *tuples*)
  - Each row gives a value for every field, from the appropriate domain.
- The *arity* of a relation is the number of fields in its schema.
- The *cardinality* of a relation is the number of rows in its table.

**Absolutely everything** in a relational database is built from relations and operations upon them.



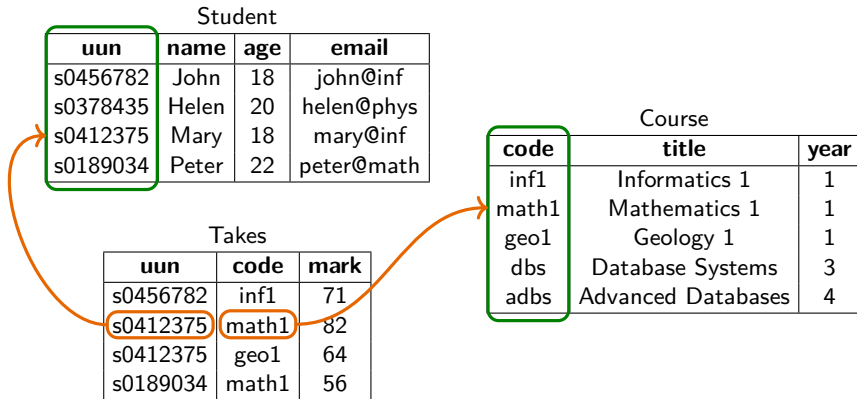
# Example

Relational databases take as fundamental the idea of a *relation*, comprising a *schema* and an *instance*.



Every relational database is a linked collection of several tables like this: often much wider, and sometimes very, very much longer.

# Some Linked Relations



- Primary Key
- Foreign Key

# SQL: Structured Query Language

- SQL is the standard language for working with relational database management systems.
- Substantial parts of SQL are **declarative**: code states what should be done, not necessarily how to do it.
- When actually querying a large database, database systems take advantage of this to plan, rearrange, and optimize the execution of queries.
- Procedural parts of SQL do contain **imperative** code to make changes to the database.
- While SQL is an international standard (ISO 9075), individual implementations have notable idiosyncrasies, and code may not be entirely portable.

There is also lots of activity in *NoSQL* databases; these are interesting, too, but we won't be doing too much on them right now. As it happens, most of what is on this slide applies to NoSQL approaches as well.

# DDL: SQL Data Definition Language

The *Data Definition Language* is a portion of SQL used to declare the schemas for relations. In particular the DDL contains the imperative command **CREATE TABLE** which sets up a fresh, empty, table with a certain schema.

For simplicity, we'll use schemas with only three types:

- **INTEGER** for integer values;
- **FLOAT** for floating point real-valued numbers;
- **VARCHAR(n)** for strings of length up to **n**.

It's conventional, although not at all universal, to write SQL keywords in **UPPER CASE**, with **lower case** or **Mixed Case** for identifiers. This isn't enforced by the language.

There are some moderately complex rules about use of quotation marks around identifiers and strings, which I'll talk about another time.

# DDL Example

```
CREATE TABLE Student (  
    uun    VARCHAR(8),  
    name   VARCHAR(20),  
    age    INTEGER,  
    email  VARCHAR(25),  
    PRIMARY KEY (uun) )
```

Student			
uun	name	age	email
s0456782	John	18	john@inf
s0378435	Helen	20	helen@phys
s0412375	Mary	18	mary@inf
s0189034	Peter	22	peter@math

The general form of this statement is

```
CREATE TABLE table-name ( <attribute declarations> <integrity constraints> )
```

with a body containing a comma-separated list of *attribute-name attribute-type* pairs and a series of statements that present constraints on the data in the table.

# DDL Example

```
CREATE TABLE Student (  
    uun    VARCHAR(8),  
    name   VARCHAR(20),  
    age    INTEGER,  
    email  VARCHAR(25),  
    PRIMARY KEY (uun) )
```

Student			
uun	name	age	email
s0456782	John	18	john@inf
s0378435	Helen	20	helen@phys
s0412375	Mary	18	mary@inf
s0189034	Peter	22	peter@math

The final line declares a *primary key constraint*, that the single field **uun** is the primary key for the **Student** table.

This states that no two rows in the **Student** table can share the same value for **uun**.

The constraint is enforced by the system: any attempt to insert a new row that duplicates an existing **uun** value will fail.

# DDL Example

```
CREATE TABLE Takes (  
  uun    VARCHAR(8),  
  code   VARCHAR(20),  
  mark   INTEGER,  
  PRIMARY KEY (uun, code),  
  FOREIGN KEY (uun) REFERENCES Student,  
  FOREIGN KEY (code) REFERENCES Course )
```

uun	code	mark
s0456782	inf1	71
s0412375	math1	82
s0412375	geo1	64
s0189034	math1	56

This relation has a *composite key* with both **uun** and **code** needed to identify a **Takes** record.

Additional *foreign key constraints* declare that **uun** will always take a value that appears in the **Student** table, and that **code** will always be from the **Course** table.

These constraints are enforced by the system and no rows can be added that do not satisfy all of them.

# From ER to Relational Models

**Entity-relationship modelling** gives a high-level conceptual design for a database system, showing what things are to be recorded and how they are connected.

The **relational model** supports a more explicit logical design, closer to implementation.

Some work is required to move between them: notice, for example, that in the ER model **Student** would have been an entity with attributes, and **Takes** a relationship between entities, while in the relational model both are implemented as tables.

In the next lecture we shall look at how to systematically transform an ER model into a relational one.